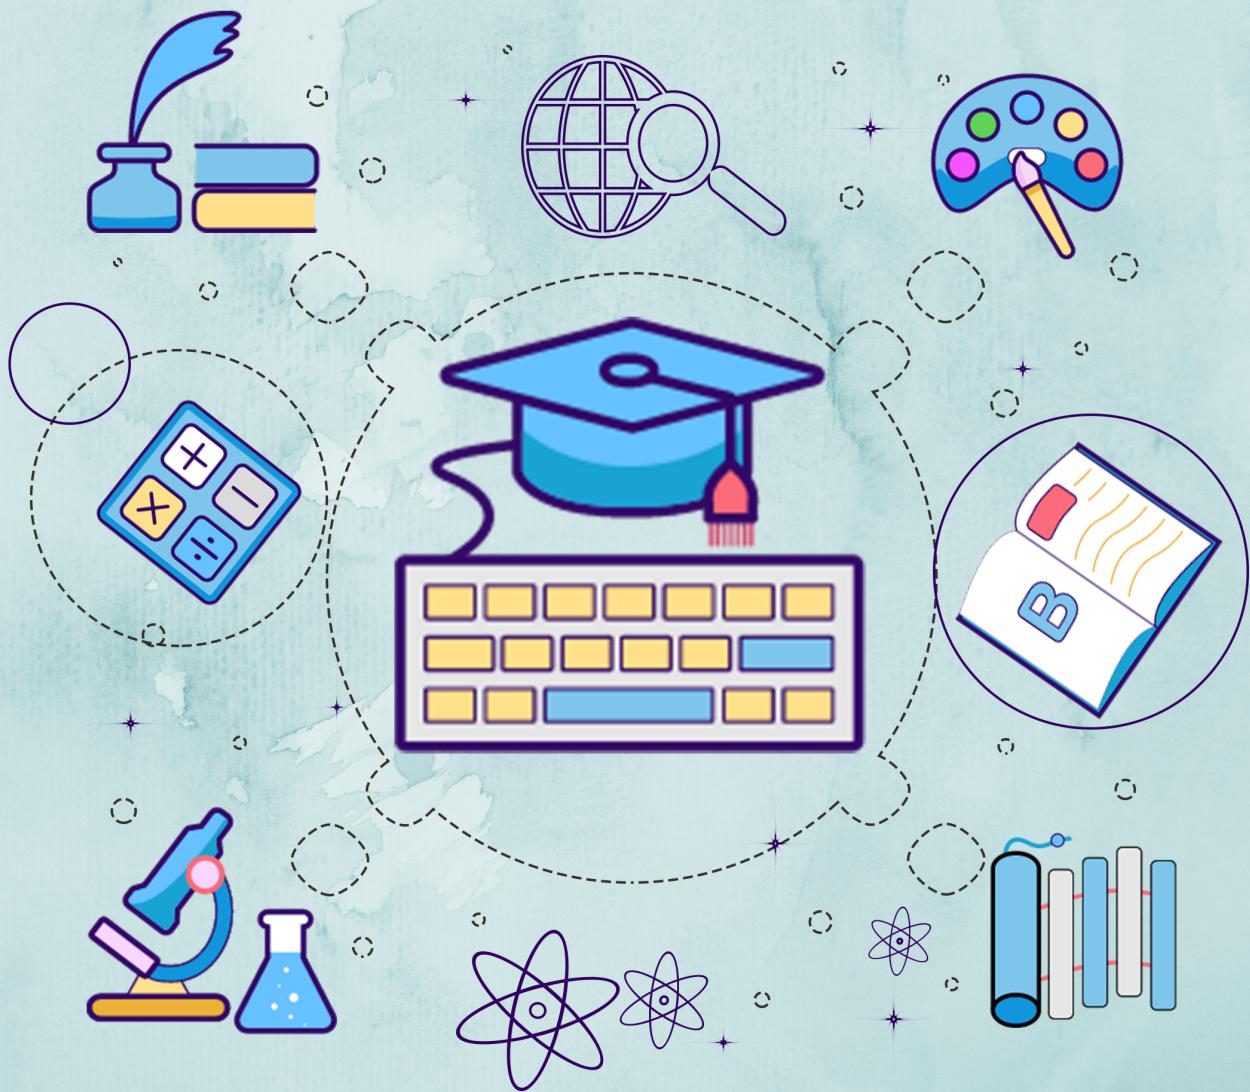


# Kerala Notes



**SYLLABUS | STUDY MATERIALS | TEXTBOOK**

**PDF | SOLVED QUESTION PAPERS**



## KTU S4 CSE PDF NOTES

# OPERATING SYSTEM (CST200)

## Module 3

### Related Link :

- KTU S4 CSE NOTES | 2019 SCHEME
- KTU S4 SYLLABUS CSE | COMPUTER SCIENCE
- KTU PREVIOUS QUESTION BANK S4 CSE SOLVED
- KTU CSE TEXTBOOKS S4 B.TECH PDF DOWNLOAD
- KTU S4 CSE NOTES | SYLLABUS | QBANK | TEXTBOOKS DOWNLOAD

## MODULE III

# PROCESS SYNCHRONISATION

## CRITICAL SECTION

- Each process has a segment of code, called a critical section in which the process may be changing common variables, updating a table, writing a file, and so on.
- The important feature of the system is that, when one process is executing in its critical section, no other process is to be allowed to execute in its critical section.
- That is, no two processes are executing in their critical sections at the same time.

### Critical Section Problem

- The critical-section problem is to design a protocol that the processes can use to cooperate. Each process must request permission to enter its critical section. The section of code implementing this request is the entry section. The critical section may be followed by an exit section. The remaining code is the remainder section.
- The general structure of a typical process  $P_i$  is shown below

```
do
{
    entry section
```

critical section

*exit section*

remainder section

} while (TRUE);

- A solution to the critical-section problem must satisfy the following three requirements:

**1. Mutual exclusion.**

- If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections.

**2. Progress.**

- If no process is executing in its critical section and some processes wish to enter their critical sections, then only those processes that are not executing in their remainder sections can participate in deciding which will enter its critical section next, and this selection cannot be postponed indefinitely.

**3 Bounded waiting.**

- There exists a bound, or limit, on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.

We assume that each process is executing at a nonzero speed. However, we can make no assumption concerning the relative speed of the  $n$  processes. At a given point in time, many kernel-mode processes may be active in the operating system. As a result, the code implementing an operating system (*kernel code*) is subject to several possible race conditions. Consider as an example a kernel data structure that maintains a list of all open files in the system. This list must be modified when a new file is opened or closed (adding the file to the list or removing it from the list). If two processes were to open files simultaneously, the separate updates to this list could result in a race condition. Other kernel data structures that are prone to possible race conditions include structures for maintaining memory allocation, for maintaining process lists, and for interrupt handling.

- Two general approaches are used to handle critical sections in operating systems:
  - (1) **pre-emptive kernels**
  - (2) **nonpreemptive kernels.**
- A pre-emptive kernel allows a process to be pre-empted while it is running in kernel mode.
- A nonpreemptive kernel does not allow a process running in kernel mode to be pre-empted. A kernel-mode process will run until it exits kernel mode, blocks, or voluntarily yields control of the CPU.

## PETERSON'S SOLUTION

- A classic software-based solution to the critical-section problem known as Peterson's solution.
- Modern computer architectures perform basic machine-language instructions, such as load and store.
- There are no guarantees that Peterson's solution will work correctly on such architectures.
- We present the solution because it provides a good algorithmic description of solving the critical-section problem
- It illustrates some of the complexities involved in designing software that addresses the requirements of mutual exclusion, progress, and bounded waiting.
- Peterson's solution is restricted to two processes that alternate execution between their critical sections and remainder sections.
- The processes are numbered  $P_0$  and  $P_1$ . For convenience, when presenting  $P_i$ , we use  $P_j$  to denote the other process; that is,  $j$  equals  $1 - i$ . Peterson's solution requires the two processes to share two data items:

```
int turn;
boolean flag[2];
```

The variable *turn* indicates whose turn it is to enter its critical section. That is, if *turn* ==  $i$ , then process  $P_i$  is allowed to execute in its critical section. The *flag* array is used to indicate if a process *is ready* to enter its critical section. For example, if *flag* [ $i$ ] is true, this value indicates that  $P_i$  is ready to enter its critical section. With an explanation of these data structures complete. The algorithm explains below.

To enter the critical section, process  $P_i$  first sets flag [i] to be true and then sets turn to the value j, thereby asserting that if the other process wishes to enter the critical section, it can do so. If both processes try to enter at the same time, turn will be set to both i and j at roughly the same time. Only one of these assignments will last; the other will occur but will be overwritten immediately.

```

Do
{
    flag [i] = TRUE;
    turn= j;
    while (flag[j] && turn j);

    critical section

    flag [i] = FALSE;

    remainder section
}
While (TRUE);

```

The eventual value of turn determines which of the two processes is allowed to enter its critical section first.

To prove that this solution is correct, it is required to show that:

- 1 Mutual exclusion is preserved.
- 2 The progress requirement is satisfied.
- 3 The bounded-waiting requirement is met.

To prove property 1, we note that each  $P_i$  enters its critical section only if either flag [j] == false or turn == i. Also note that, if both processes can be executing in their critical sections at the same time, then flag [0] == flag [1] == true. These two observations imply that  $P_0$  and  $P_1$  could not have successfully executed their while statements at about the same time, since the value of turn can be either 0 or 1 but cannot be both. Hence, one of the processes -say,  $P_i$  - must have successfully executed the while statement, whereas  $P_j$  had to execute at least one additional statement ("turn== j"). However, at that time, flag [j] == true and turn == j, and this condition will persist as long as  $P_i$  is in its critical section; as a result, mutual exclusion is preserved. To prove properties 2 and 3, we note that a process  $P_i$  can be prevented from entering the critical section only if it is stuck in the while loop with the condition flag [j] == true and turn== j; this loop is the only one possible. If  $P_i$  is not ready to enter the critical section, then flag [j] == false, and  $P_i$  can enter its critical section. If  $P_j$  has set flag [j] to true and is also executing in its while statement, then either turn == i or turn == j. If turn == i, then  $P_i$  will enter the critical section. If turn== j, then  $P_i$  will enter the critical section. However, once  $P_i$  exits its critical section, it will reset flag [j] to false, allowing  $P_j$  to enter its critical section. If  $P_i$  resets flag [j] to true, it must also set turn to i. Thus, since  $P_i$  does not change the value of the variable turn while executing the while statement,  $P_i$  will enter the critical section (progress) after at most one entry by  $P_j$  (bounded waiting).

# SYNCHRONIZATION

## SYNCHRONIZATION HARDWARE

It can generally state that any solution to the critical-section problem requires a simple tool-a **lock**. Race conditions are prevented by requiring that critical regions be protected by locks. That is, a process must acquire a lock before entering a critical section; it releases the lock when it exits the critical section.

The critical-section problem could be solved simply in a single-processor environment if it could prevent interrupts from occurring while a shared variable was being modified. In this way, it could be sure that the current sequence of instructions would be allowed to execute in order without preemption. No other instructions would be run, so no unexpected modifications could be made to the shared variable. This is often the approach taken by nonpreemptive kernels.

Disabling interrupts on a multiprocessor can be time consuming, since the message is passed to all the processors. This message passing delays entry into each critical section, and system efficiency decreases.

Many modern computer systems therefore provide special hardware instructions that allow us either to test and modify the content of a word or to swap the contents of two words atomically.

➤ Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

The definition of the test and set() instruction

➤ Solution:

```
do{
while (test_and_set(&lock))
    ; /* do nothing */
    /* critical section */
lock = false;
    /* remainder section */
} while (true);
```

Mutual-exclusion implementation with test and set().

The compare and swap() instruction, operates on three operands.

➤ Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

The definition of the compare and swap() instruction.

➤ Solution:

```
do{
while (compare_and_swap(&lock, 0, 1) != 0)
    ; /* do nothing */
    /* critical section */
```

```

lock = 0;
/* remainder section */
} while (true);

```

Mutual-exclusion implementation with the compare and swap() instruction

### **Mutex Locks**

- Simplest tools to solve the critical-section problem is **mutex lock**.
- A process must acquire the lock before entering a critical section; it releases the lock when it exits the critical section.
- The acquire() function acquires the lock, and the release() function releases the lock
- A mutex lock has a boolean variable **available** whose value indicates if the lock is available or not. If the lock is available, a call to acquire() succeeds, and the lock is then considered unavailable. A process that attempts to acquire an unavailable lock is blocked until the lock is released.
- Solution to the critical-section problem using mutex locks.

```

acquire() {
    while (!available)
        ; /* busy wait */
    available = false;;
}

```

```

do
{
    acquire lock

```

critical section

*release lock*

remainder section

```

} while (TRUE);

```

- The definition of release() is

```

release() {
    available = true;
}

```

- Calls to either acquire() or release() must be performed atomically.
- The main disadvantage of the implementation given here is that it requires busy waiting. While a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the call to acquire(). In fact, this type of mutex lock is also called a **spinlock** because the process “spins” while waiting for the lock to become available.
-

## SEMAPHORE

The hardware-based solutions to the critical-section problem presented are complicated for application programmers to use. To overcome this difficulty a synchronization tool called a semaphore is used. **A semaphore S is an integer variable that, apart from initialization, is accessed only through two standard atomic operations: wait () and signal ().** The wait () operation was originally termed P (from the Dutch *proberen*, "to test"); signal() was originally called V (from *verhogen*, "to increment"). The definition of wait () is as follows:

```
wait(S)
{
    while S <= 0
        // busy waiting
        S--;
}

}
```

The definition of signal() is as follows:

```
signal(S)
{
    S++;
}
```

All modifications to the integer value of the semaphore in the wait () and signal() operations must be executed indivisibly. That is, when one process modifies the semaphore value, no other process can simultaneously modify that same semaphore value. In addition, in the case of wait (S), the testing of the integer value of S ( $S \leq 0$ ), as well as its possible modification ( $S--$ ), must be executed without interruption.

### Semaphore Usage

Operating systems often distinguish between counting and binary semaphores.

- The value of a **counting semaphore** can range over an unrestricted domain.
- The value of a **binary semaphore** can range only between 0 and 1.

Counting semaphores can be used to control access to a given resource consisting of a finite number of instances. The semaphore is initialized to the number of resources available. Each process that wishes to use a resource performs a wait() operation on the semaphore (thereby decrementing the count). When a process releases a resource, it performs a signal() operation (incrementing the count). When the count for the semaphore goes to 0, all resources are being used. After that, processes that wish to use a resource will block until the count becomes greater than 0.

- Consider  $P_1$  and  $P_2$  that require  $S_1$  to happen before  $S_2$   
 Create a semaphore “synch” initialized to 0

P1:

```
S1;
signal(synch);
```

P2:

```
wait(synch);
S2;
```

## Semaphore Implementation

When a process executes the wait() operation and finds that the semaphore value is not positive, it must wait. However, rather than engaging in busy waiting, the process can block itself. The block operation places a process into a waiting queue associated with the semaphore, and the state of the process is switched to the waiting state. Then control is transferred to the CPU scheduler, which selects another process to execute.

A process that is blocked, waiting on a semaphore S, should be restarted when some other process executes a signal() operation. The process is restarted by a wakeup() operation, which changes the process from the waiting state to the ready state. The process is then placed in the ready queue.

- a semaphore can be defined as follows:

```
typedef struct{

    int value;

    struct process *list;

} semaphore;
```

- The wait() semaphore operation can be defined as

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}
```

- the signal() semaphore operation can be defined as

```
signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

The block() operation suspends the process that invokes it. The wakeup(P) operation resumes the execution of a blocked process P.

If a semaphore value is negative, its magnitude is the number of processes waiting on that semaphore. This fact results from switching the order of the decrement and the test in the implementation of the wait() operation.

## Deadlocks and Starvation

The implementation of a semaphore with a waiting queue may result in a situation where two or more processes are waiting indefinitely for an event that can be caused only by one of the waiting processes. When such a state is reached, these processes are said to be deadlocked.

- Let S and Q be two semaphores initialized to 1

$P_0$	$P_1$
wait(S);	wait(Q);
wait(Q);	wait(S);
...	...
signal(S);	signal(Q);
signal(Q);	signal(S);

Suppose that  $P_0$  executes wait(S) and then  $P_1$  executes wait(Q). When  $P_0$  executes wait(Q), it must wait until  $P_1$  executes signal(Q). Similarly, when  $P_1$  executes wait(S), it must wait until  $P_0$  executes signal(S). Since these signal() operations cannot be executed,  $P_0$  and  $P_1$  are deadlocked.

Another problem related to deadlocks is indefinite blocking or starvation, a situation in which processes wait indefinitely within the semaphore.

## Priority Inversion

- Scheduling problem when lower-priority process holds a lock needed by higher-priority process is called priority inversion.
- Solution
  - to have only two priorities
  - implement a priority-inheritance protocol- According to this protocol, all processes that are accessing resources needed by a higher-priority process inherit the higher priority until they are finished with the resources in question. When they are finished, their priorities revert to their original values.

## Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- A monitor type is an ADT(*Abstract data type*) that includes a set of programmer defined operations that are provided with mutual exclusion within the monitor. The monitor type also declares the variables whose values define the state of an instance of that type, along with the bodies of functions that operate on those variables.

monitor monitor name

```
{  
/* shared variable declarations */  
function P1(...) {  
...  
}  
function P2(...) {  
...  
}  
...  
function Pn(...) {  
...  
}  
initialization code (...) {  
...  
}
```

Syntax of a monitor.

- A function defined within a monitor can access only those variables declared locally within the monitor and its formal parameters.
- The local variables of a monitor can be accessed by only the local functions.
- The monitor construct ensures that only one process at a time is active within the monitor.

Schematic view of a monitor

- A monitor consists of a **mutex (lock)** object and **condition variables**. A **condition variable** is basically a container of threads that are waiting for a certain condition. Monitors provide a mechanism for threads to temporarily give up exclusive access in order to wait for some condition to be met, before regaining exclusive access and resuming their task.
- The only operations that can be invoked on a condition variable are `wait()` and `signal()`. The operation `x.wait();` means that the process invoking this operation is suspended until another process invokes. The `x.signal();` operation resumes exactly one suspended process. If no process is suspended, then the `signal()` operation has no effect.

#### Monitor with condition variables.

- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`,
  - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
  - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
  - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition

## Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem

## Bounded-Buffer Problem

- The producer and consumer processes share the following data structures:

```
int n;  
  
semaphore mutex = 1;  
  
semaphore empty = n;  
  
semaphore full = 0
```

- The pool consists of n buffers, each capable of holding one item. The mutex semaphore provides mutual exclusion for accesses to the buffer pool and is initialized to the value 1. The empty and full semaphores count the number of empty and full buffers. The semaphore empty is initialized to the value n; the semaphore full is initialized to the value 0.
- The producer producing full buffers for the consumer or as the consumer producing empty buffers for the producer.
- The structure of the producer process.

```
do {  
  
    ...  
    /* produce an item in next_produced */  
  
    ...  
    wait(empty);  
    wait(mutex);  
  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
}  
} while (true);
```

- The structure of the consumer process

```

do {

    wait(full);

    wait(mutex);

    ...

    /* remove an item from buffer to next_consumed */

    ...

    signal(mutex);

    signal(empty);

    ...

    /* consume the item in next consumed */

    ...

} while (true);

```

### **Readers-Writers Problem**

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do *not* perform any updates
  - Writers – can both read and write
- Problem –
  - allow multiple readers to read at the same time
  - Only one single writer can access the shared data at the same time
- To ensure that these difficulties do not arise, we require that the writers have exclusive access to the shared database while writing to the database. This synchronization problem is referred to as the readers-writers problem
- The readers-writers problem has several variations,
  - **The first readers-writers problem** requires that no reader be kept waiting unless a writer has already obtained permission to use the shared object. In other words, no reader should wait for other readers to finish simply because a writer is waiting
  - **The second readers –writers problem** requires that, once a writer is ready, that writer perform its write as soon as possible. In other words, if a writer is waiting to access the object, no new readers may start reading.
- A solution to either problem may result in starvation.

➤ Shared Data

```
semaphore rw_mutex = 1;
semaphore mutex = 1;
int read_count = 0;
```

- The semaphores mutex and rw mutex are initialized to 1; read count is initialized to 0. The semaphore rw mutex is common to both reader and writer processes.
- The mutex semaphore is used to ensure mutual exclusion when the variable read count is updated. The read count variable keeps track of how many processes are currently reading the object. The semaphore rw mutex functions as a mutual exclusion semaphore for the writers. It is also used by the first or last reader that enters or exits the critical section. It is not used by readers who enter or exit while other readers are in their critical sections.

➤ The structure of a writer process

```
do {
    wait(rw_mutex);

    ...
    /* writing is performed */

    ...
    signal(rw_mutex);

} while (true);
```

➤ The structure of a reader process

```
do {
    do {
        wait(mutex);
        read_count++;
        if (read_count == 1)
            wait(rw_mutex);
        signal(mutex);

        ...
        /* reading is performed */

        ...
        wait(mutex);
        read_count--;
        if (read_count == 0)

        signal(rw_mutex);

        signal(mutex);
    } while (true);
}
```

- The readers-writers problem and its solutions have been generalized to provide reader-writer locks on some systems. Acquiring a reader-writer lock requires specifying the mode of the lock: either read or write access.
- When a process wishes only to read shared data, it requests the reader-writer lock in read mode. A process wishing to modify the shared data must request the lock in write mode. Multiple processes are permitted to concurrently acquire a reader-writer lock in read mode, but only one process may acquire the lock for writing, as exclusive access is required for writers.

### The Dining-Philosophers Problem

The situation of the dining philosophers

- Philosophers spend their lives alternating thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
  - Need both to eat, then release both when done
- The shared data are

```
semaphore chopstick[5];
```
- Where all the elements of chopstick are initialized to 1.
- The structure of Philosopher  $i$ :

```
do {
    wait(chopstick[i]);
    wait(chopstick[(i+1) % 5]);
    ...
    /* eat for awhile */
    ...
    signal(chopstick[i]);
    signal(chopstick[(i+1) % 5]);
    ...
    /* think for awhile */
    ... } while (true);
```

## DEADLOCK



KeralaNotes

Process require resource, if not available moves to waiting state. The resources that are requested are held by waiting process.

### System Model

Set of Resources → Logical resources

[files, semaphore]

→ Physical resources

[CPU, memory, printer etc]

A process must request a resource before using it and must release the resource after using it.

The no. of resources requested may not exceed the total no. of resources available in the system.

### Normal mode of operation:

1) Request: If the request cannot be granted immediately ie, the resource is being used by another process, then the requesting process must wait until it can acquire the resource.

2) Use: The process can operate on the resource if granted, ie, if the resource is a printer, the process can print on the printer.

3) Release: The process releases the resources.

### Deadlock Characteristics

#### 4) Necessary Conditions:

Deadlock condition hold if all the four conditions hold simultaneously:

(i) Mutual Exclusion

- Atleast one resource must be held in non-shareable mode i.e., if only one process at a time can use the resource.
- If another process requests that resource, the requesting process must be delayed until the resource has been released.

(ii) Hold and wait

- A process must be holding atleast one resource & waiting to acquire additional resources that are currently being held by other processes.

(iii) No Preemption

- Resources cannot be preempted i.e., a resource can be released only voluntarily by the process holding it, after that process has completed its task.

(iv) Circular Wait

- A set  $\{P_0, P_1, P_2, \dots, P_n\}$  of waiting process must exist such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2, \dots, P_{n-1}$  is waiting for a resource that is held by  $P_n$  and  $P_n$  is waiting for a resource that is held by  $P_0$ .



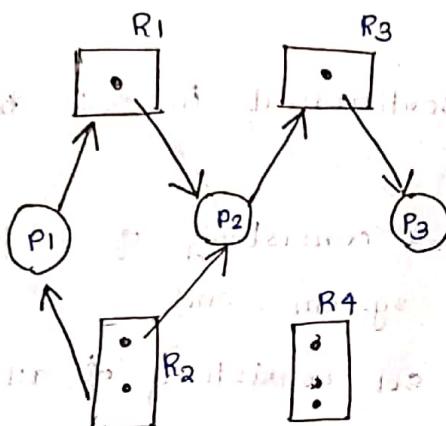
## 2) Resource Allocation Graph

- Deadlocks can be described in terms of directed graph called a system resource-allocation graph.
- The graph consists of a set of vertices  $V$  & a set of edges  $E$ .
- The set of vertices  $V$  is partitioned into different types of nodes:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the active process in the system.
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system.
- A directed edge from process  $P_i$  to resource type  $R_j$  is denoted by  $P_i \rightarrow R_j$ .  
ie, The Process  $P_i$  is requested an instance of resource type  $R_j$  & is currently waiting for that resource. It is called request edge.
- A directed edge from resource type  $R_j$  to process  $P_i$  is denoted by  $R_j \rightarrow P_i$ .  
ie, An instance of resource type  $R_j$  has been allocated to process  $P_i$ . It is called Assignment edge.
- Process as circle, resource as square.
- Since request type  $R_j$  has more than one instance, we represent each such instance as a dot within square.
- When Process  $P_i$  requests an instance of resource type  $R_j$ , a request edge is inserted in the resource-allocation graph. When this request can be fulfilled the request edge is transformed to assignment edge. When the process no longer needs access to the resource it releases the resource. As a result assignment edge is deleted.

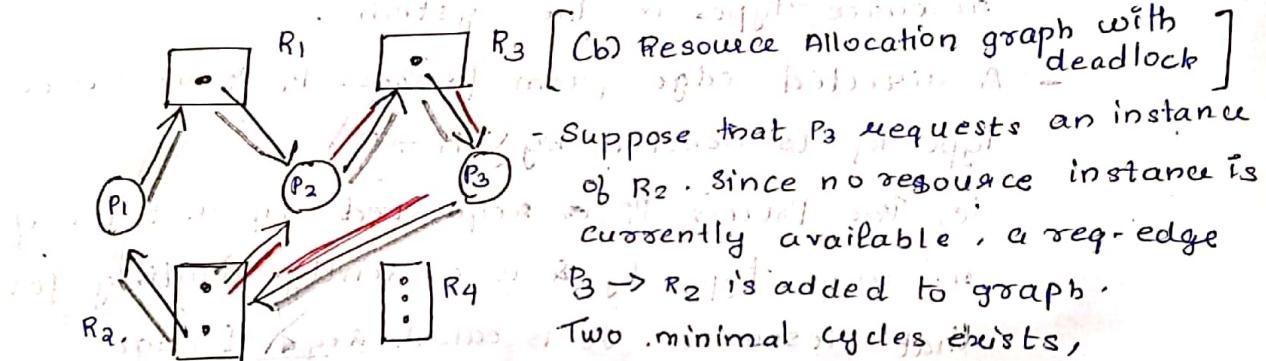
- If the graph contains no cycle, then no process in the system is deadlocked.

- If the graph does contain a cycle, then a deadlock may exist.

- If each resource type has exactly one instance, then a cycle implies a deadlock has occurred.



[a] Resource Allocation graph



[b] Resource Allocation graph with deadlock

- Suppose that P3 requests an instance of R2. Since no resource instance is currently available, a req-edge  $P_3 \rightarrow R_2$  is added to graph.

Two minimal cycles exists,

$$P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$$

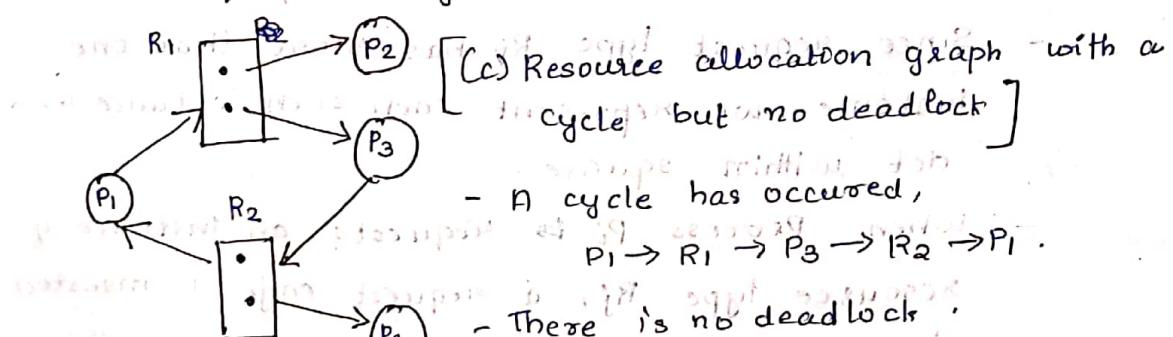
$$P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$$

✓ P1, P2, P3 are deadlocked.

✓ P2 is waiting for R3 which is held by P3.

✓ P3 is waiting for either P1 or P2 to release R2.

✓ P1 is waiting for P2 to release R1.



[c] Resource allocation graph with a cycle, but no deadlock

- A cycle has occurred,

$$P_1 \rightarrow R_1 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$$

- There is no deadlock.

- P4 may release its instance of R2.

That resource can then be

allocated to P3, breaking the cycle.

## Methods of Handling Deadlocks

We can deal with the deadlock problem in one of three ways:

1. Can use a protocol to prevent or avoid deadlock, ensuring that the system will never enter a deadlock state.
2. Can allow the system to enter a deadlock state, detect it & recovery.
3. Can ignore the problem altogether, & pretend the deadlocks never occur in the system.

To ensure that deadlocks never occur, the system can use either a deadlock prevention, nor a deadlock avoidance scheme.

- Deadlock Prevention is a set of methods for ensuring that atleast one of the necessary conditions cannot hold.
- Deadlock Avoidance requires that the OS to give additional information about which resources a process will request & use.

### Deadlock Prevention

Ensure that atleast one condition among 4 doesn't hold for the deadlock to happen.

1. Mutual Exclusion.

- Mutual exclusion condition must hold for non-shareable resources.

- Eg: A printer cannot be simultaneously shared by several processes.

- Shareable resources do not require mutually exclusive access & thus cannot be involved in a deadlock.

**For More Study Materials : <https://www.keralanotes.com/>**

- If several processes attempt to open a read-only file at the same time, they can be granted simultaneous access to the file.

- A process never needs to wait for a sharable resource.

## 2. Hold and Wait.

- Whenever a process requests a resource, it does not hold any other resources.

- One protocol that can be used requires each process to request & be allocated all its resources before it begins execution.

- An alternative protocol allows a process to request resource only when the process has none. A process may request some resources & use them. Before it can request any additional resources, however it must release all the resources that it is currently allocated.

We consider a process that copies data from tape drive to a disk files, sort the disk files & then prints the results to a printer. If all the resources must be requested at the beginning of the process, then the process must initially request the tape drive, disk file & printer. It will hold the printer for its entire execution, even though it needs the printer only at the end.

The second method allows the process to request initially only the tape drive & disk file. It copies from the tape drive to disk, then releases both the tape drive & disk file. The process must then again request the disk file & printer. After copying the disk file to the printer, it releases these 2 resources & terminates.

### 3. No Preemption

$$P_i \rightarrow R_j \rightarrow P_j$$

-  $P_i$  wait until  $P_j$  releases  $R_j$



- If a process is holding some resources & requests another resources that cannot be immediately allocated it, ie, the process must wait, then all resources currently being held are preempted.
- The preempted resources are added to the list of resources for which the process is waiting. The process will be restarted only when it can regain its old resources, as well as the new ones that is requesting.
- Alternatively if a process requests some resources, we first check whether they are available. If they are, we allocate them.  $P_i \rightarrow R_j \rightarrow P_j \rightarrow R_k$ .  $P_i \rightarrow R_j$  then  $R_j \rightarrow P_j$ .  
- If they are not available, we check whether they are allocated to some other process that is waiting for additional resource. If so, we preempt the desired resources from the waiting process & allocate them to the requesting process.  
- If the resources are not either available or held by a waiting process, the requesting process must wait while it is waiting, some of its resources may be preempted but only if another process requests them. A process can be restarted only when it is allocated the new resources it is requesting & recovers any resource that were preempted while it was waiting.
- One way to ensure that this condition never holds is to impose a total ordering of all resource types & to require that each process requests resources in an increasing order of enumeration.

### 4. Circular Wait

One way to ensure that this condition never holds is to impose a total ordering of all resource types & to require that each process requests resources in an increasing order of enumeration.

**For More Study Materials : <https://www.keralanotes.com/>**

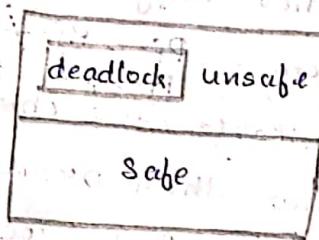
A Function  $F: R \rightarrow N$ , where  $N$  is the set of natural numbers.

$$F(R_0) < F(R_1) < \dots < F(R_n) < F(R_0)$$

## Deadlock Avoidance:

- For Avoiding deadlock it is to require additional information about how resources have to be requested.
- The simplest and most useful model requires that each process declare the maximum no. of resources of each type that it may need.

### 4) Safe State:



- A state is safe, if the system can allocate resources to each process in some order & still avoid a deadlock.
- A system is in a safe state only if there is a safe sequence.
- A sequence of processes  $\langle P_1, P_2, \dots, P_n \rangle$  is a safe sequence, if for each  $P_i$ , the resources that  $P_i$  ~~still~~ requests can be satisfied by the currently available resources plus the resources held by all  $P_j$ .
- If the resource that process  $P_i$  needs are not immediately available, then  $P_i$  can wait until all  $P_j$  has finished. When they have finished,  $P_i$  can obtain all of its needed resources, complete its task, return its allocated resources & terminate. When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resource & so on.

- A safe state is not a deadlock state.
- A deadlock state is an unsafe state.
- Not all unsafe states are deadlocks.

### Problem:

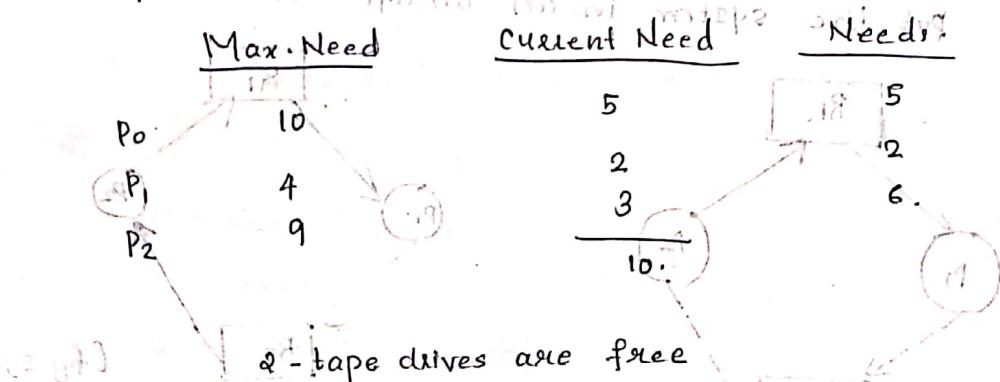
Total : 12 magnetic tape drives.

	<u>Max. Need</u>	<u>Current Need</u>	<u>Needs</u>
P <sub>0</sub>	10	5	5
P <sub>1</sub>	4	2	2
P <sub>2</sub>	9	2	7

all 6 units 3-tape drives are free. & all processes are in ready queue. Then  $\langle P_1, P_0, P_2 \rangle$  -safe state. First  $P_1$  is allocated & then when it releases it have 5 available.

Then  $P_0$  can get all its tape drives & return them, now it may have 10 available & finally  $P_2$  could get all 10 & it may return as 12. Thus a deadlock occurs.

Suppose if  $P_2$  requests & is allocated 4 tape drives.



2-tape drives are free.

Here,  $P_1$  can only be allocated all its tape drives.

when it returns, it has 4. Since  $P_0$  needs 5 & it cannot satisfy the need of  $P_0$  or  $P_2$ . Thus a deadlock occurs.

deadlock arises due to starvation of processes and

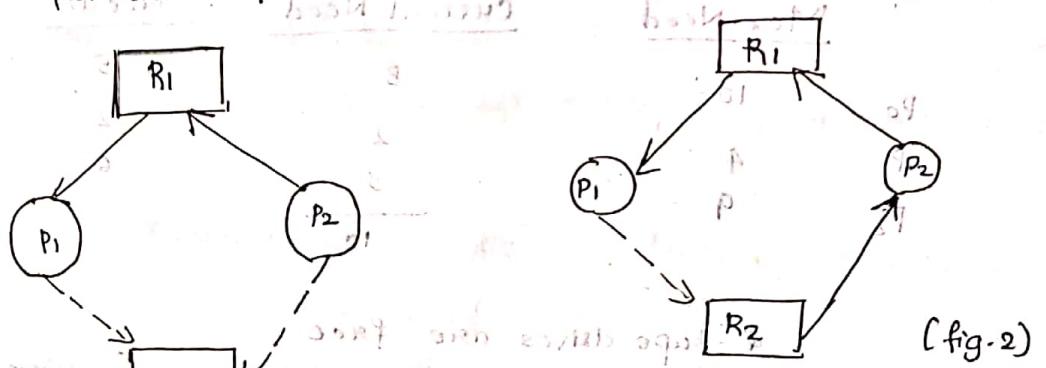
abuse of memory and insufficient resources.

in such a condition, it is better to terminate the process.

## 2) Resource Allocation Graph.

A claim edge  $P_i \rightarrow R_j$  indicates that process  $P_i$  requests resource  $R_j$  at some time in future. This edge resembles a request edge in direction, but is represented by dashed lines. When process  $P_i$  requests resource  $R_j$ , the claim edge  $P_i \rightarrow R_j$  is converted to a request edge. Similarly, when a resource  $R_j$  is released by  $P_j$ , the assignment edge  $R_j \rightarrow P_i$  is converted to a claim edge  $P_i \rightarrow R_j$ . Suppose that process  $P_i$  requests resource  $R_j$ . The resource can be granted only if converting the request edge  $P_i \rightarrow R_j$  to an assignment edge  $R_j \rightarrow P_i$  does not result in the formation of a cycle in resource allocation graph.

If no cycle exists, then the allocation of resource will leave the system in a safe state. If a cycle is found, then the allocation will put the system in an unsafe state.



[Fig-1] - Suppose that  $P_2$  requests  $R_2$ . Although  $R_2$  is currently free, we cannot allocate it to  $P_2$ , since this action will create a cycle in the graph.

[Fig-2] - A cycle indicates that the system is unsafe.

If  $P_1$  requests  $R_2$  &  $P_2$  requests  $R_1$ , then a deadlock will occur.

## Banker's Algorithm

↳ Available [j] = k

- k instance of resource type  $R_j$  available.

↳ Max [i, j] = ~~Max~~.

-  $P_i$  require atmost k instance of resource type  $R_j$ .

↳ Allocation [i, j] = k.

- Process  $P_i$  is currently allocated k instance of resource type  $R_j$ .

↳ Need [i, j]

- Process  $P_i$  need k more instance of resource type  $R_j$  to complete its task.

$$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j].$$

## Safety Algorithm

1) Let work & finish be vectors of length m & n resp.

Initialize work = Available &

Finish [i] := false, for  $i=1, 2, \dots, n$ .

2) Find an i such that both -

(a) Finish [i] = false

(b) Need[i]  $\leq$  work.

If no such i exists then go to step 4.

3) work = work + Allocation<sub>i</sub>.

Finish [i] := true

Go to step 2.

4) If Finish [i] = true for all i, then the system is in a safe state.

## Resource Request Algorithm

- 1) If  $\underline{\text{Request}} \leq \text{Need}_i$ , go to step 2.  
Otherwise raise an error condition, since the process has executed its maximum claim.
  - 2) If  $\underline{\text{Request}}_i \leq \text{Available}$ , go to step 3.  
Otherwise  $P_i$  must wait, since the resources are not available.
  - 3) Have the system pretend to have allocated the requested resource to process  $P_i$  by

Available := Available - Request;

$$\checkmark \text{Allocation} := \text{Allocation} + \text{Request}$$

$$\checkmark \text{Need}_i := \text{Need}_i^{\circ} - \text{Request}_i$$

## Problem

4)

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	0	1	0	7	5	3	3	3	2
P <sub>1</sub>	2	0	0	3	2	2			
P <sub>2</sub>	3	0	2	9	0	2			
P <sub>3</sub>	2	1	1	2	2	2			
P <sub>4</sub>	0	0	2	4	3	3			

Soln:

## Safety Algorithm

Need [Max - Allocation] Need  $\leq$  Avail

A B C      ① P<sub>0</sub> 743  $\leq$  332 X

P<sub>0</sub> 7 4 3      ② P<sub>1</sub> 122  $\leq$  332 ✓

P<sub>1</sub> 1 2 2      Avail = 332 + 200 = 532

P<sub>2</sub> 6 0 0      ③ P<sub>2</sub> 600  $\leq$  532 X

P<sub>3</sub> 0 1 1      ④ P<sub>3</sub> 011  $\leq$  532 ✓

P<sub>4</sub> 4 3 1      Avail = 532 + 211 = 743

⑤ P<sub>4</sub> 431  $\leq$  743 ✓

Avail = 743 + 002 = 745

⑥ P<sub>0</sub> 743  $\leq$  745 ✓

Avail = 745 + 010 = 755

⑦ P<sub>2</sub> 600  $\leq$  755 ✓

∴ Avail = 755 + 302 = 1057

∴ A B C

10 5 7

∴ {P<sub>1</sub>, P<sub>3</sub>, P<sub>4</sub>, P<sub>0</sub>, P<sub>2</sub>} is safe.

Request of P<sub>1</sub> comes to 102  $\leq$  122 ✓

Request of P<sub>1</sub> comes to 102  $\leq$  332 ✓

Avail := 332 - 102 = 230 [Avail := Avail - Req]

Alloc := 200 + 102 = 302 [Alloc := Alloc + Req]

Need = 122 - 102 = 020 [Need := Need - Req]

	<u>Alloc</u>	<u>Max</u>	<u>Avail</u>	<u>Need</u>
	302	322	230	020

For More Study Materials : <https://www.keralanotes.com/>

Part

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>			<u>Need</u>		
	A	B	C	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	0	1	0	7	5	3	2	3	0	7	4	3
P <sub>1</sub>	3	0	2	8	2	2				0	8	0
P <sub>2</sub>	3	0	2	9	0	2				6	0	0
P <sub>3</sub>	2	1	1	2	2	2				0	1	1
P <sub>4</sub>	0	0	2	4	3	3				4	3	1
				8	2	7						

De

① P<sub>0</sub>. 743  $\leq$  230 X.

② P<sub>1</sub>. 020  $\leq$  230 ✓

$$230 + 302 = 532.$$

③ P<sub>2</sub>. 600  $\leq$  532 X.

④ P<sub>3</sub>. 011  $\leq$  532 ✓

$$532 + 211 = 743.$$

⑤ P<sub>4</sub>. 431  $\leq$  743 ✓

$$743 + 002 = 745.$$

⑥ P<sub>0</sub>. 743  $\leq$  745 ✓

$$745 + 010 = 755.$$

⑦ P<sub>2</sub>. 600  $\leq$  755 ✓

$$755 + 302 = 1057 \quad \therefore \langle P_1, P_3, P_4, P_0, P_2 \rangle \text{ is safe}$$

Ab Req. P<sub>4</sub> (3 3 0) & Req. P<sub>0</sub> (0 2 0) comes.

Req. Algorithm.

① 330  $\leq$  431 ✓ 330  $\leq$  230 X. Avail = 230 - 020 = 210

② 020  $\leq$  743 ✓ 020  $\leq$  230 ✓ All = 010 + 020 = 030. Need = 743 - 020 = 723.

	<u>Allocation</u>			<u>Need</u>			<u>Avail</u>			
	A	B	C	A	B	C	A	B	C	
P <sub>0</sub>	0	3	0	7	2	3	2	1	0	② P <sub>1</sub> . 020 $\leq$ 210 X
P <sub>1</sub>	3	0	2	0	2	0				③ P <sub>2</sub> . 600 $\leq$ 210 X
P <sub>2</sub>	3	0	2	6	0	0				④ P <sub>3</sub> . 011 $\leq$ 210 X
P <sub>3</sub>	2	1	1	0	1	1				⑤ P <sub>4</sub> . 431 $\leq$ 210 X
P <sub>4</sub>	0	0	2	4	3	1				Undo the changes, since it is unsafe

so. So change as in Part 1

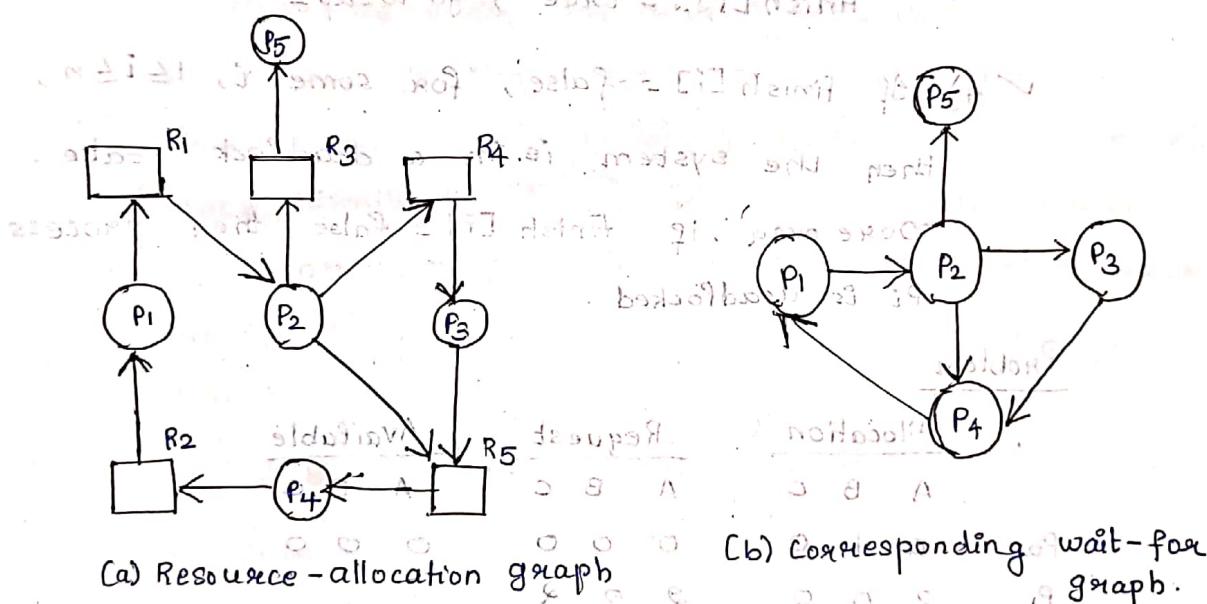
## Deadlock Detection

Determine whether deadlock has occurred.

### 1) Single Instance of Each Resource Type.

If all the resources have only a single instance, then we can define a deadlock detection algorithm called a wait-for graph.

An Edge from  $P_i$  to  $P_j$  in a wait-for graph implies that process  $P_i$  is waiting for process  $P_j$  to release a resource that  $P_i$  needs. An edge  $P_i \rightarrow P_j$  exists in a wait-for graph if & only if the corresponding resource-allocation graph contains two edges  $P_i \rightarrow R_q$  &  $R_q \rightarrow P_j$  for some resource  $R_q$ .



### 2.) Several Instance of a Resource Type.

The wait-for graph scheme is not applicable to a resource allocation system with multiple instances of each resource type.

### Deadlock detection algorithm

1) Let work & finish be vectors of length  $n$  respectively.

Initialize work<sub>i</sub> = available, for  $i=0, 1, 2, \dots, n$ .

If Allocation  $\leq$  work, then Finish [i] = false

Otherwise

Finish [i] = true.

2) Find an index  $i$  such that both:

(a) Finish [i] = false.

✓ (b) Request<sub>i</sub>  $\leq$  work<sub>i</sub>.

If no such  $i$  exists, go to step 4.

3) work<sub>i</sub> = work + Allocation.

Finish [i] := true, go to step 2.

✓ 4) If Finish [i] = false, for some  $i$ ,  $1 \leq i \leq n$ ,

then the system is in a deadlock state.

Moreover, if Finish [i] = false, then process  $P_i$  is deadlocked.

### Problem -

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	0	1	0	0	0	0	0	0	0
P <sub>1</sub>	2	0	0	2	0	2	0	0	0
P <sub>2</sub>	3	0	3	0	0	0	0	0	0
P <sub>3</sub>	2	1	1	1	0	0	0	0	0
P <sub>4</sub>	0	0	2	0	0	2	0	0	0
P <sub>0</sub>	$1 \leq \text{avail}$			$0 \leq \text{req}$			$0 \leq \text{work}$		
P <sub>0</sub>	$000 \leq 000 \checkmark$			$000 \leq 000 \checkmark$			$000 \leq 000 \checkmark$		
P <sub>1</sub>	$202 \leq 010 \times$			$202 \leq 010 \times$			$000 + 010 = 010$		
P <sub>2</sub>	$000 \leq 010 \checkmark$			$010 + 303 = 313$			$000 + 010 = 010$		
P <sub>3</sub>	$100 \leq 213 \checkmark$			$213 + 211 = 524$			$000 + 010 = 010$		
P <sub>4</sub>	$002 \leq 524 \checkmark$			$524 + 002 = 526$			$000 + 010 = 010$		
				$526 + 200 = 726$					

## Recovery from Deadlock

### 1) Process Termination

- ✓ Abort all deadlocked processes.

- This method clearly will break the deadlock cycle, but at a great expense.

- ✓ Abort one process at a time until deadlock cycle is eliminated.

- This method incurs considerable overhead, since after each process is aborted in a deadlock detection algorithm must be invoked to determine whether any processes are still deadlocked.

Aborting a process may not be easy as the process was in the midst of updating a file, terminating it will leave that file in an incorrect state. Similarly, if the process was in the midst of printing data on the printer, the system must reset the printer to a correct state before printing the next job.

Criteria for process selection :忙iest

- ✓ what the priority of the process is.
- ✓ How long process was computed.
- ✓ How many & what type of resource used.
- ✓ How many more resources needed.
- ✓ How many process will need to be terminated.

## Q) Resource Preemption



Issues regarding resource preemption

### ✓ Selecting a victim

- Which resources & which processes are to be preempted? As in process termination, we must determine the order of preemption to minimize cost of interrupting a process & breaking it.

### ✓ Roll back

- If we preempt a resource from a process, then it cannot continue with its normal execution; it is missing some needed resource. We must roll back the process to some safer state & restart it from that state.

### ✓ Starvation

In a system where victim selection is based primarily on cost factors, it may happen that the same process is always picked as a victim. As a result, this process never completes its designated tasks, a starvation situation arises.

Starvation is a situation where a process is denied a resource it needs for a long time.

Causes of starvation: -

1. Shortage of required resources.

2. Shortage of time available for execution.

3. Shortage of memory space.

For More Study Materials : <https://www.keralanotes.com/>